

Keras -- MLPs on MNIST

In [0]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [0]:

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [4]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 2s 0us/step

In [5]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [7]:

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [7]:

```
# An example data point
print(X_train[0])
```

[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	255
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	154
170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	0
0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	82	82
82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	253
253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154	154	154
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241	241
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	187	187
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	253
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	39	148	229													

In [0]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normlizing
print(X_train[0])
```

[illegible]

[illegible]

[illegible]

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:
```

```

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [0]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [13]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66:
The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541:
The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432:
: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

MLP + ReLU + ADAM

In [14]:

```
model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(72, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4409:
: The name tf.random_normal is deprecated. Please use tf.random.normal instead.

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 364)	285740
dense_3 (Dense)	(None, 72)	26280
dense_4 (Dense)	(None, 10)	730
Total params: 312,750		
Trainable params: 312,750		
Non-trainable params: 0		

None

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576:
: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250:
add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 107us/step - loss: 0.2509 - acc: 0.9257 - val_loss: 0.1234 - val_acc: 0.9620

Epoch 2/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0985 - acc: 0.9710 - val_loss: 0.0895 - val_acc: 0.9728

Epoch 3/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0618 - acc: 0.9816 - val_loss: 0.0827 - val_acc: 0.9757

Epoch 4/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0424 - acc: 0.9871 - val_loss: 0.0909 - val_acc: 0.9735

Epoch 5/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0311 - acc: 0.9904 - val_loss: 0.0701 - val_acc: 0.9802

```

0701 - val_acc: 0.9803
Epoch 6/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0223 - acc: 0.9932 - val_loss: 0.
0770 - val_acc: 0.9773
Epoch 7/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0191 - acc: 0.9939 - val_loss: 0.
0859 - val_acc: 0.9771
Epoch 8/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0133 - acc: 0.9959 - val_loss: 0.
0859 - val_acc: 0.9789
Epoch 9/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0133 - acc: 0.9956 - val_loss: 0.
0984 - val_acc: 0.9740
Epoch 10/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0125 - acc: 0.9958 - val_loss: 0.
0779 - val_acc: 0.9807
Epoch 11/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0105 - acc: 0.9966 - val_loss: 0.
0813 - val_acc: 0.9803
Epoch 12/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0120 - acc: 0.9958 - val_loss: 0.
0822 - val_acc: 0.9790
Epoch 13/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0082 - acc: 0.9972 - val_loss: 0.
0981 - val_acc: 0.9770
Epoch 14/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0078 - acc: 0.9976 - val_loss: 0.
1054 - val_acc: 0.9765
Epoch 15/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0063 - acc: 0.9977 - val_loss: 0.
0899 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0091 - acc: 0.9971 - val_loss: 0.
1151 - val_acc: 0.9776
Epoch 17/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0103 - acc: 0.9965 - val_loss: 0.
1113 - val_acc: 0.9785
Epoch 18/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0068 - acc: 0.9976 - val_loss: 0.
0904 - val_acc: 0.9817
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0035 - acc: 0.9990 - val_loss: 0.
0891 - val_acc: 0.9813
Epoch 20/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0100 - acc: 0.9967 - val_loss: 0.
0927 - val_acc: 0.9810

```

In [33]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))
#print(x)

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']

```

```
#print (ty)
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09268554465530374
Test accuracy: 0.981



In [35]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [14]:

```
# Multilayer perceptron
```



```

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ 
.
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
    mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(72, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4409: The name tf.random_normal is deprecated. Please use tf.random.normal instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 364)	285740
batch_normalization_1 (Batch Normalization)	(None, 364)	1456
dense_3 (Dense)	(None, 72)	26280
batch_normalization_2 (Batch Normalization)	(None, 72)	288
dense_4 (Dense)	(None, 10)	730
Total params: 314,494		
Trainable params: 313,622		
Non-trainable params: 872		

In [15]:

```

model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 158us/step - loss: 0.2165 - acc: 0.9376 - val_loss: 0.1199 - val_acc: 0.9646

Epoch 2/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0828 - acc: 0.9759 - val_loss: 0.0924 - val_acc: 0.9718

Epoch 3/20

```

60000/60000 [=====] - 5s 87us/step - loss: 0.0549 - acc: 0.9833 - val_loss: 0.
0867 - val_acc: 0.9731
Epoch 4/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0380 - acc: 0.9889 - val_loss: 0.
0847 - val_acc: 0.9739
Epoch 5/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0293 - acc: 0.9910 - val_loss: 0.
0862 - val_acc: 0.9750
Epoch 6/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0230 - acc: 0.9930 - val_loss: 0.
0749 - val_acc: 0.9774
Epoch 7/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0188 - acc: 0.9942 - val_loss: 0.
0807 - val_acc: 0.9763
Epoch 8/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0153 - acc: 0.9954 - val_loss: 0.
0813 - val_acc: 0.9789
Epoch 9/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0167 - acc: 0.9945 - val_loss: 0.
0835 - val_acc: 0.9756
Epoch 10/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0151 - acc: 0.9954 - val_loss: 0.
0931 - val_acc: 0.9735
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0115 - acc: 0.9961 - val_loss: 0.
0823 - val_acc: 0.9774
Epoch 12/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0112 - acc: 0.9964 - val_loss: 0.
0856 - val_acc: 0.9769
Epoch 13/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0112 - acc: 0.9965 - val_loss: 0.
0815 - val_acc: 0.9785
Epoch 14/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0079 - acc: 0.9974 - val_loss: 0.
0840 - val_acc: 0.9782
Epoch 15/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0085 - acc: 0.9973 - val_loss: 0.
0887 - val_acc: 0.9786
Epoch 16/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0078 - acc: 0.9977 - val_loss: 0.
0880 - val_acc: 0.9780
Epoch 17/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0089 - acc: 0.9973 - val_loss: 0.
0959 - val_acc: 0.9759
Epoch 18/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0092 - acc: 0.9971 - val_loss: 0.
0871 - val_acc: 0.9774
Epoch 19/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0066 - acc: 0.9977 - val_loss: 0.
0918 - val_acc: 0.9783
Epoch 20/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0080 - acc: 0.9974 - val_loss: 0.
0905 - val_acc: 0.9767

```

In [16]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss


```

```
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09051545914988383

Test accuracy: 0.9767



In [17]:


```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + Dropout + AdamOptimizer

In [18]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(72, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed
=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733
: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed
in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 364)	285740
batch_normalization_3 (Batch Normalization)	(None, 364)	1456
dropout_1 (Dropout)	(None, 364)	0
dense_6 (Dense)	(None, 72)	26280
batch_normalization_4 (Batch Normalization)	(None, 72)	288
dropout_2 (Dropout)	(None, 72)	0
dense_7 (Dense)	(None, 10)	730
Total params: 314,494		
Trainable params: 313,622		
Non-trainable params: 872		

In [19]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 101us/step - loss: 0.5634 - acc: 0.8287 - val_loss: 0.1879 - val_acc: 0.9407

Epoch 2/20

60000/60000 [=====] - 5s 85us/step - loss: 0.2963 - acc: 0.9125 - val_loss: 0.1416 - val_acc: 0.9558

Epoch 3/20

60000/60000 [=====] - 5s 84us/step - loss: 0.2462 - acc: 0.9276 - val_loss: 0.1202 - val_acc: 0.9638

Epoch 4/20

60000/60000 [=====] - 5s 84us/step - loss: 0.2090 - acc: 0.9383 - val_loss: 0.1052 - val_acc: 0.9660

Epoch 5/20

60000/60000 [=====] - 5s 87us/step - loss: 0.1861 - acc: 0.9445 - val_loss: 0.1003 - val_acc: 0.9688

```

10000 - val_acc: 0.9700
Epoch 6/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1717 - acc: 0.9499 - val_loss: 0.
0941 - val_acc: 0.9720
Epoch 7/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1558 - acc: 0.9544 - val_loss: 0.
0899 - val_acc: 0.9725
Epoch 8/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1437 - acc: 0.9575 - val_loss: 0.
0835 - val_acc: 0.9728
Epoch 9/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1335 - acc: 0.9595 - val_loss: 0.
0785 - val_acc: 0.9769
Epoch 10/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1249 - acc: 0.9628 - val_loss: 0.
0778 - val_acc: 0.9764
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.1209 - acc: 0.9645 - val_loss: 0.
0813 - val_acc: 0.9773
Epoch 12/20
60000/60000 [=====] - 5s 81us/step - loss: 0.1189 - acc: 0.9647 - val_loss: 0.
0749 - val_acc: 0.9774
Epoch 13/20
60000/60000 [=====] - 5s 81us/step - loss: 0.1117 - acc: 0.9665 - val_loss: 0.
0788 - val_acc: 0.9774
Epoch 14/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1086 - acc: 0.9674 - val_loss: 0.
0734 - val_acc: 0.9792
Epoch 15/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1040 - acc: 0.9686 - val_loss: 0.
0693 - val_acc: 0.9802
Epoch 16/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1003 - acc: 0.9696 - val_loss: 0.
0706 - val_acc: 0.9795
Epoch 17/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0961 - acc: 0.9708 - val_loss: 0.
0704 - val_acc: 0.9810
Epoch 18/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0921 - acc: 0.9727 - val_loss: 0.
0691 - val_acc: 0.9816
Epoch 19/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0876 - acc: 0.9735 - val_loss: 0.
0647 - val_acc: 0.9806
Epoch 20/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0861 - acc: 0.9730 - val_loss: 0.
0624 - val_acc: 0.9818

```

In [20]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0624060621908051

Test accuracy: 0.9818



In [21]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 Layer MLP+ReLU+Adam

In [22]:

```
model_relu = Sequential()
model_relu.add(Dense(250, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, se
```

```

ed=None)) )
model_relu.add(Dense(85, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 250)	196250
dense_9 (Dense)	(None, 100)	25100
dense_10 (Dense)	(None, 85)	8585
dense_11 (Dense)	(None, 10)	860

Total params: 230,795
 Trainable params: 230,795
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 62us/step - loss: 0.2668 - acc: 0.9218 - val_loss: 0.1340 - val_acc: 0.9586

Epoch 2/20

60000/60000 [=====] - 3s 53us/step - loss: 0.1021 - acc: 0.9693 - val_loss: 0.0955 - val_acc: 0.9698

Epoch 3/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0665 - acc: 0.9791 - val_loss: 0.0776 - val_acc: 0.9765

Epoch 4/20

60000/60000 [=====] - 3s 52us/step - loss: 0.0463 - acc: 0.9858 - val_loss: 0.0721 - val_acc: 0.9773

Epoch 5/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0371 - acc: 0.9882 - val_loss: 0.0875 - val_acc: 0.9743

Epoch 6/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0297 - acc: 0.9901 - val_loss: 0.0786 - val_acc: 0.9767

Epoch 7/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0236 - acc: 0.9923 - val_loss: 0.0724 - val_acc: 0.9791

Epoch 8/20

60000/60000 [=====] - 3s 52us/step - loss: 0.0208 - acc: 0.9931 - val_loss: 0.0805 - val_acc: 0.9791

Epoch 9/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0190 - acc: 0.9938 - val_loss: 0.0845 - val_acc: 0.9782

Epoch 10/20

60000/60000 [=====] - 3s 52us/step - loss: 0.0164 - acc: 0.9944 - val_loss: 0.0820 - val_acc: 0.9798

Epoch 11/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0119 - acc: 0.9960 - val_loss: 0.1087 - val_acc: 0.9746

Epoch 12/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0162 - acc: 0.9948 - val_loss: 0.0803 - val_acc: 0.9818

Epoch 13/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0126 - acc: 0.9961 - val_loss: 0.0965 - val_acc: 0.9790

Epoch 14/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0124 - acc: 0.9960 - val_loss: 0.0963 - val_acc: 0.9781

Epoch 15/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0130 - acc: 0.9960 - val_loss: 0.0880 - val_acc: 0.9798

Epoch 16/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0086 - acc: 0.9972 - val_loss: 0.0880 - val_acc: 0.9818

```

60000/60000 [-----] - 3s 34us/step - loss: 0.0066 - acc: 0.9972 - val_loss: 0.
1445 - val_acc: 0.9678
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0109 - acc: 0.9963 - val_loss: 0.
0896 - val_acc: 0.9798
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0067 - acc: 0.9978 - val_loss: 0.
0978 - val_acc: 0.9794
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0123 - acc: 0.9960 - val_loss: 0.
1034 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0076 - acc: 0.9978 - val_loss: 0.
0878 - val_acc: 0.9827

```

In [23]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))
#print(x)

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
#print(ty)
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.087810692615682

Test accuracy: 0.9827



In [26]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

```



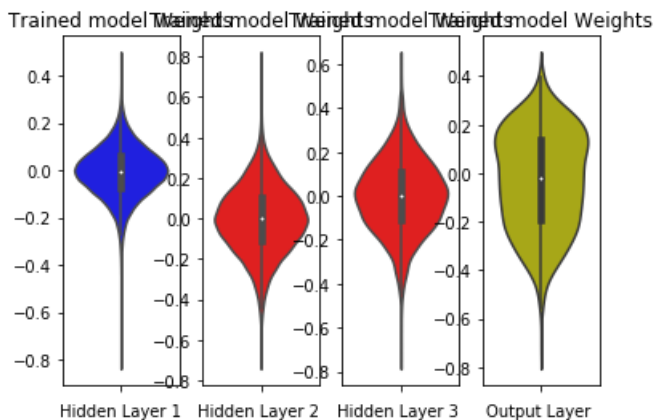
```
out_w = w_attr[0].flatten().reshape(1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')
```

```
plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 Layer MLP+Batch Normalization+ReLU+Adam

In [27]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ 
.
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h3 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(250, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
    mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(85, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 250)	196250
batch_normalization_5 (Batch Normalization)	(None, 250)	1000
dense_13 (Dense)	(None, 100)	25100
batch_normalization_6 (Batch Normalization)	(None, 100)	400
dense_14 (Dense)	(None, 85)	8585
batch_normalization_7 (Batch Normalization)	(None, 85)	340
dense_15 (Dense)	(None, 10)	860

Total params: 232,535

Trainable params: 231,665

Non-trainable params: 870

In [28]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 118us/step - loss: 0.2506 - acc: 0.9276 - val_loss: 0.1278 - val_acc: 0.9606

Epoch 2/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0959 - acc: 0.9715 - val_loss: 0.1056 - val_acc: 0.9669

Epoch 3/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0645 - acc: 0.9806 - val_loss: 0.0965 - val_acc: 0.9693

Epoch 4/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0471 - acc: 0.9851 - val_loss: 0.0907 - val_acc: 0.9718

Epoch 5/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0364 - acc: 0.9891 - val_loss: 0.0841 - val_acc: 0.9746

Epoch 6/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0302 - acc: 0.9905 - val_loss: 0.0857 - val_acc: 0.9755

Epoch 7/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0265 - acc: 0.9912 - val_loss: 0.0807 - val_acc: 0.9755

Epoch 8/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0240 - acc: 0.9923 - val_loss: 0.0807 - val_acc: 0.9759

Epoch 9/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0197 - acc: 0.9935 - val_loss: 0.0865 - val_acc: 0.9750

Epoch 10/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0153 - acc: 0.9952 - val_loss: 0.0768 - val_acc: 0.9784

Epoch 11/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0132 - acc: 0.9959 - val_loss: 0.0849 - val_acc: 0.9762

Epoch 12/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0138 - acc: 0.9952 - val_loss: 0.0880 - val_acc: 0.9776

Epoch 13/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0118 - acc: 0.9960 - val_loss: 0.0923 - val_acc: 0.9762

Epoch 14/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0130 - acc: 0.9956 - val_loss: 0.0923 - val_acc: 0.9762

```

0.0858 - val_acc: 0.9783
Epoch 15/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0100 - acc: 0.9967 - val_loss: 0.
0947 - val_acc: 0.9774
Epoch 16/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0097 - acc: 0.9967 - val_loss: 0.
0912 - val_acc: 0.9781
Epoch 17/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0111 - acc: 0.9963 - val_loss: 0
.0892 - val_acc: 0.9783
Epoch 18/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0109 - acc: 0.9965 - val_loss: 0
.0922 - val_acc: 0.9764
Epoch 19/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0085 - acc: 0.9971 - val_loss: 0.
1031 - val_acc: 0.9772
Epoch 20/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0061 - acc: 0.9982 - val_loss: 0
.0889 - val_acc: 0.9800

```

In [29]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08889246554333222

Test accuracy: 0.98



In [30]:

```

w_after = model_batch.get_weights()
h1 w = w_after[0].flatten().reshape(-1,1)

```

```

h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Trained model Weights")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



3 Layer MLP+ReLU+Dropout+Adam

In [31]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(250, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(85, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

```

```
model_drop.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 250)	196250
batch_normalization_8 (Batch Normalization)	(None, 250)	1000
dropout_3 (Dropout)	(None, 250)	0
dense_17 (Dense)	(None, 100)	25100
batch_normalization_9 (Batch Normalization)	(None, 100)	400
dropout_4 (Dropout)	(None, 100)	0
dense_18 (Dense)	(None, 85)	8585
batch_normalization_10 (Batch Normalization)	(None, 85)	340
dropout_5 (Dropout)	(None, 85)	0
dense_19 (Dense)	(None, 10)	860
Total params: 232,535		
Trainable params: 231,665		
Non-trainable params: 870		

In [32]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 128us/step - loss: 1.0290 - acc: 0.6731 - val_loss: 0.2667 - val_acc: 0.9234

Epoch 2/20

60000/60000 [=====] - 6s 105us/step - loss: 0.4735 - acc: 0.8611 - val_loss: 0.1968 - val_acc: 0.9387

Epoch 3/20

60000/60000 [=====] - 6s 106us/step - loss: 0.3619 - acc: 0.8965 - val_loss: 0.1666 - val_acc: 0.9508

Epoch 4/20

60000/60000 [=====] - 6s 105us/step - loss: 0.3176 - acc: 0.9091 - val_loss: 0.1467 - val_acc: 0.9554

Epoch 5/20

60000/60000 [=====] - 6s 106us/step - loss: 0.2827 - acc: 0.9201 - val_loss: 0.1331 - val_acc: 0.9597

Epoch 6/20

60000/60000 [=====] - 6s 107us/step - loss: 0.2512 - acc: 0.9293 - val_loss: 0.1291 - val_acc: 0.9605

Epoch 7/20

60000/60000 [=====] - 6s 105us/step - loss: 0.2323 - acc: 0.9352 - val_loss: 0.1202 - val_acc: 0.9652

Epoch 8/20

60000/60000 [=====] - 6s 104us/step - loss: 0.2164 - acc: 0.9398 - val_loss: 0.1096 - val_acc: 0.9664

Epoch 9/20

60000/60000 [=====] - 6s 103us/step - loss: 0.2037 - acc: 0.9433 - val_loss: 0.1117 - val_acc: 0.9674

Epoch 10/20

60000/60000 [=====] - 6s 105us/step - loss: 0.1891 - acc: 0.9478 - val_loss: 0.1063 - val_acc: 0.9697

Epoch 11/20

60000/60000 [=====] - 6s 106us/step - loss: 0.1822 - acc: 0.9496 - val_loss: 0.0971 - val_acc: 0.9712

Epoch 12/20

60000/60000 [=====] - 6s 106us/step - loss: 0.1717 - acc: 0.9513 - val_loss: 0.0944 - val_acc: 0.9734

```

Epoch 13/20
60000/60000 [=====] - 6s 104us/step - loss: 0.1608 - acc: 0.9560 - val_loss: 0
.0919 - val_acc: 0.9739
Epoch 14/20
60000/60000 [=====] - 6s 103us/step - loss: 0.1577 - acc: 0.9553 - val_loss: 0
.0886 - val_acc: 0.9754
Epoch 15/20
60000/60000 [=====] - 6s 105us/step - loss: 0.1487 - acc: 0.9583 - val_loss: 0
.0853 - val_acc: 0.9754
Epoch 16/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1448 - acc: 0.9597 - val_loss: 0
.0872 - val_acc: 0.9758
Epoch 17/20
60000/60000 [=====] - 6s 103us/step - loss: 0.1369 - acc: 0.9622 - val_loss: 0
.0914 - val_acc: 0.9757
Epoch 18/20
60000/60000 [=====] - 6s 105us/step - loss: 0.1405 - acc: 0.9615 - val_loss: 0
.0871 - val_acc: 0.9762
Epoch 19/20
60000/60000 [=====] - 6s 104us/step - loss: 0.1317 - acc: 0.9621 - val_loss: 0
.0839 - val_acc: 0.9758
Epoch 20/20
60000/60000 [=====] - 6s 104us/step - loss: 0.1272 - acc: 0.9645 - val_loss: 0
.0853 - val_acc: 0.9760

```

In [33]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08526910331975669

Test accuracy: 0.976



In [34]:

```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5 Layer MLP+ReLU+Adam

In [35]:

```
model_relu = Sequential()
model_relu.add(Dense(235, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, se
ed=None)))
model_relu.add(Dense(65, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, see
d=None)))
model_relu.add(Dense(50, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, see
d=None)))
model_relu.add(Dense(35, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, see
d=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validatio
n_data=(X_test, Y_test))
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 235)	184475
dense_21 (Dense)	(None, 102)	24072
dense_22 (Dense)	(None, 65)	6695
dense_23 (Dense)	(None, 50)	3300
dense_24 (Dense)	(None, 35)	1785
dense_25 (Dense)	(None, 10)	360

Total params: 220,687

Trainable params: 220,687

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 79us/step - loss: 0.3140 - acc: 0.9043 - val_loss: 0.1389 - val_acc: 0.9569

Epoch 2/20

60000/60000 [=====] - 4s 63us/step - loss: 0.1132 - acc: 0.9657 - val_loss: 0.1291 - val_acc: 0.9598

Epoch 3/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0774 - acc: 0.9762 - val_loss: 0.1000 - val_acc: 0.9690

Epoch 4/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0598 - acc: 0.9810 - val_loss: 0.0857 - val_acc: 0.9748

Epoch 5/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0437 - acc: 0.9863 - val_loss: 0.0832 - val_acc: 0.9769

Epoch 6/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0346 - acc: 0.9893 - val_loss: 0.0851 - val_acc: 0.9767

Epoch 7/20

60000/60000 [=====] - 4s 64us/step - loss: 0.0315 - acc: 0.9896 - val_loss: 0.0810 - val_acc: 0.9772

Epoch 8/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0242 - acc: 0.9918 - val_loss: 0.0936 - val_acc: 0.9745

Epoch 9/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0237 - acc: 0.9920 - val_loss: 0.1073 - val_acc: 0.9735

Epoch 10/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0221 - acc: 0.9927 - val_loss: 0.0905 - val_acc: 0.9774

Epoch 11/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0198 - acc: 0.9936 - val_loss: 0.0849 - val_acc: 0.9778

Epoch 12/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0183 - acc: 0.9938 - val_loss: 0.1086 - val_acc: 0.9759

Epoch 13/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0176 - acc: 0.9945 - val_loss: 0.0952 - val_acc: 0.9766

Epoch 14/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0149 - acc: 0.9951 - val_loss: 0.1125 - val_acc: 0.9751

Epoch 15/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0131 - acc: 0.9959 - val_loss: 0.0983 - val_acc: 0.9779

Epoch 16/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0132 - acc: 0.9959 - val_loss: 0.0997 - val_acc: 0.9783

Epoch 17/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0130 - acc: 0.9959 - val_loss: 0.1009 - val_acc: 0.9797

Epoch 18/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0125 - acc: 0.9964 - val_loss: 0.0887 - val_acc: 0.9814


```

0.007 - val_acc: 0.9765
Epoch 19/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0081 - acc: 0.9975 - val_loss: 0.
1159 - val_acc: 0.9765
Epoch 20/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0130 - acc: 0.9959 - val_loss: 0.
1187 - val_acc: 0.9765

```

In [36]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.11873800104274924

Test accuracy: 0.9765



In [38]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(v=h1_w, color='b')

```

```

plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')


plt.subplot(1, 6, 3)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='y')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 6, 4)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='b')
plt.xlabel('Hidden Layer 4')

plt.subplot(1, 6, 5)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer')
plt.show()

```



5 Layer MLP+ Batch normalization_ReLu_Adam

In [39]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
.
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(235, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
    mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(65, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

```

```

model_batch.add(Dense(50, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(35, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 235)	184475
batch_normalization_11 (Batch Normalization)	(None, 235)	940
dense_27 (Dense)	(None, 102)	24072
batch_normalization_12 (Batch Normalization)	(None, 102)	408
dense_28 (Dense)	(None, 65)	6695
batch_normalization_13 (Batch Normalization)	(None, 65)	260
dense_29 (Dense)	(None, 50)	3300
batch_normalization_14 (Batch Normalization)	(None, 50)	200
dense_30 (Dense)	(None, 35)	1785
batch_normalization_15 (Batch Normalization)	(None, 35)	140
dense_31 (Dense)	(None, 10)	360
Total params: 222,635		
Trainable params: 221,661		
Non-trainable params: 974		

In [40]:

```

model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 11s 183us/step - loss: 0.3773 - acc: 0.8937 - val_loss: 0.1553 - val_acc: 0.9558
Epoch 2/20
60000/60000 [=====] - 9s 143us/step - loss: 0.1254 - acc: 0.9630 - val_loss: 0.1210 - val_acc: 0.9621
Epoch 3/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0879 - acc: 0.9733 - val_loss: 0.1143 - val_acc: 0.9656
Epoch 4/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0680 - acc: 0.9791 - val_loss: 0.1035 - val_acc: 0.9705
Epoch 5/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0527 - acc: 0.9833 - val_loss: 0.1121 - val_acc: 0.9672
Epoch 6/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0434 - acc: 0.9863 - val_loss: 0.1037 - val_acc: 0.9702
Epoch 7/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0364 - acc: 0.9882 - val_loss: 0.1156 - val_acc: 0.9676

```

```

.1100 - val_acc: 0.9735
Epoch 8/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0346 - acc: 0.9883 - val_loss: 0
.0918 - val_acc: 0.9739
Epoch 9/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0305 - acc: 0.9895 - val_loss: 0
.0995 - val_acc: 0.9735
Epoch 10/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0276 - acc: 0.9914 - val_loss: 0
.1009 - val_acc: 0.9726
Epoch 11/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0238 - acc: 0.9922 - val_loss: 0
.0880 - val_acc: 0.9767
Epoch 12/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0200 - acc: 0.9933 - val_loss: 0
.0920 - val_acc: 0.9772
Epoch 13/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0201 - acc: 0.9935 - val_loss: 0
.0994 - val_acc: 0.9717
Epoch 14/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0189 - acc: 0.9934 - val_loss: 0
.0869 - val_acc: 0.9772
Epoch 15/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0155 - acc: 0.9952 - val_loss: 0
.0950 - val_acc: 0.9758
Epoch 16/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0164 - acc: 0.9945 - val_loss: 0
.0927 - val_acc: 0.9765
Epoch 17/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0158 - acc: 0.9946 - val_loss: 0
.0972 - val_acc: 0.9758
Epoch 18/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0151 - acc: 0.9948 - val_loss: 0
.0891 - val_acc: 0.9789
Epoch 19/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0108 - acc: 0.9964 - val_loss: 0
.0893 - val_acc: 0.9773
Epoch 20/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0129 - acc: 0.9956 - val_loss: 0
.1032 - val_acc: 0.9754

```

In [41]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10316739412156312

Test accuracy: 0.9754





In [42]:

```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='b')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5 MLP+Dropout+ReLU+Adam

In [43]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(235, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(65, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(50, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(35, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_32 (Dense)	(None, 235)	184475
batch_normalization_16 (Batch Normalization)	(None, 235)	940
dropout_6 (Dropout)	(None, 235)	0
dense_33 (Dense)	(None, 102)	24072
batch_normalization_17 (Batch Normalization)	(None, 102)	408
dropout_7 (Dropout)	(None, 102)	0
dense_34 (Dense)	(None, 65)	6695
batch_normalization_18 (Batch Normalization)	(None, 65)	260
dropout_8 (Dropout)	(None, 65)	0
dense_35 (Dense)	(None, 50)	3300
batch_normalization_19 (Batch Normalization)	(None, 50)	200
dropout_9 (Dropout)	(None, 50)	0
dense_36 (Dense)	(None, 35)	1785
batch_normalization_20 (Batch Normalization)	(None, 35)	140

dropout_10 (Dropout)	(None, 35)	0
dense_37 (Dense)	(None, 10)	360
=====		
Total params: 222,635		
Trainable params: 221,661		
Non-trainable params: 974		
=====		

In [44]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 12s 199us/step - loss: 2.1567 - acc: 0.2681 - val_loss: 0.9967 - val_acc: 0.7442
Epoch 2/20
60000/60000 [=====] - 9s 152us/step - loss: 1.3065 - acc: 0.5476 - val_loss: 0.5086 - val_acc: 0.8687
Epoch 3/20
60000/60000 [=====] - 9s 152us/step - loss: 0.9505 - acc: 0.6814 - val_loss: 0.3757 - val_acc: 0.8998
Epoch 4/20
60000/60000 [=====] - 9s 152us/step - loss: 0.7723 - acc: 0.7492 - val_loss: 0.3148 - val_acc: 0.9044
Epoch 5/20
60000/60000 [=====] - 9s 152us/step - loss: 0.6576 - acc: 0.7985 - val_loss: 0.2575 - val_acc: 0.9332
Epoch 6/20
60000/60000 [=====] - 9s 151us/step - loss: 0.5789 - acc: 0.8294 - val_loss: 0.2122 - val_acc: 0.9481
Epoch 7/20
60000/60000 [=====] - 9s 150us/step - loss: 0.5112 - acc: 0.8555 - val_loss: 0.1928 - val_acc: 0.9506
Epoch 8/20
60000/60000 [=====] - 9s 150us/step - loss: 0.4727 - acc: 0.8710 - val_loss: 0.1762 - val_acc: 0.9554
Epoch 9/20
60000/60000 [=====] - 9s 152us/step - loss: 0.4358 - acc: 0.8865 - val_loss: 0.1687 - val_acc: 0.9554
Epoch 10/20
60000/60000 [=====] - 9s 152us/step - loss: 0.4043 - acc: 0.8961 - val_loss: 0.1642 - val_acc: 0.9589
Epoch 11/20
60000/60000 [=====] - 9s 150us/step - loss: 0.3823 - acc: 0.9035 - val_loss: 0.1555 - val_acc: 0.9621
Epoch 12/20
60000/60000 [=====] - 9s 151us/step - loss: 0.3582 - acc: 0.9100 - val_loss: 0.1546 - val_acc: 0.9620
Epoch 13/20
60000/60000 [=====] - 9s 152us/step - loss: 0.3387 - acc: 0.9183 - val_loss: 0.1523 - val_acc: 0.9631
Epoch 14/20
60000/60000 [=====] - 9s 151us/step - loss: 0.3162 - acc: 0.9233 - val_loss: 0.1443 - val_acc: 0.9649
Epoch 15/20
60000/60000 [=====] - 9s 152us/step - loss: 0.3127 - acc: 0.9246 - val_loss: 0.1361 - val_acc: 0.9681
Epoch 16/20
60000/60000 [=====] - 9s 150us/step - loss: 0.2990 - acc: 0.9295 - val_loss: 0.1344 - val_acc: 0.9675
Epoch 17/20
60000/60000 [=====] - 10s 161us/step - loss: 0.2919 - acc: 0.9312 - val_loss: 0.1393 - val_acc: 0.9681
Epoch 18/20
60000/60000 [=====] - 9s 153us/step - loss: 0.2806 - acc: 0.9355 - val_loss: 0.1357 - val_acc: 0.9678
Epoch 19/20
60000/60000 [=====] - 9s 152us/step - loss: 0.2688 - acc: 0.9363 - val_loss: 0.1293 - val_acc: 0.9715
Epoch 20/20
```

```
Epoch 20/20
60000/60000 [=====] - 9s 152us/step - loss: 0.2613 - acc: 0.9396 - val_loss: 0.1308 - val_acc: 0.9706
```

In [45]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.13083188042640687

Test accuracy: 0.9706



In [46]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
#plt.title("Trained model Weights")
ax = sns.violinplot(v=h2_w, color='r')
```



```

plt.subplot(1, 6, 2)
plt.xlabel('Hidden Layer 2 ')

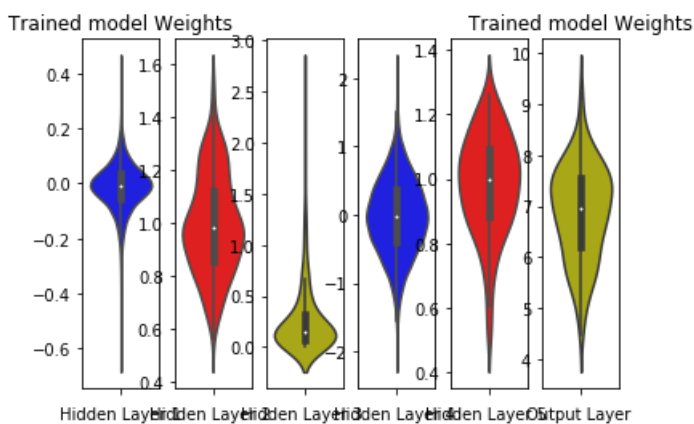
plt.subplot(1, 6, 3)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='b')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
#plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



In [49]:

```

from prettytable import PrettyTable
x = PrettyTable(["Feature", "Accuracy(%)", "Loss(%)"])
x.add_row(["ReLU+Adam", 98.1, 9.2])
x.add_row(["ReLU+Batch+Adam", 97.6, 9.05])
x.add_row(["ReLU+Dropout+Adam", 98.18, 6.24])
print("2 Layer MLP Results")
print(x)

y = PrettyTable(["Feature", "Accuracy(%)", "Loss(%)"])
y.add_row(["ReLU+Adam", 98.27, 8.78])
y.add_row(["ReLU+Batch+Adam", 98, 8.88])
y.add_row(["ReLU+Dropout+Adam", 97.6, 8.52])
print("3 Layer MLP Results")
print(y)

z = PrettyTable(["Feature", "Accuracy(%)", "Loss(%)"])
z.add_row(["ReLU+Adam", 97.65, 11.87])
z.add_row(["ReLU+Batch+Adam", 97.54, 10.31])
z.add_row(["ReLU+Dropout+Adam", 97.06, 13.08])
print("5 Layer MLP Results")
print(z)

```

2 Layer MLP Results

Feature	Accuracy(%)	Loss(%)
ReLU+Adam	98.1	9.2
ReLU+Batch+Adam	97.6	9.05
ReLU+Dropout+Adam	98.18	6.24

ReLU+Adam	98.1	9.2
ReLU+Batch+Adam	97.6	9.05
ReLU+Dropout+Adam	98.18	6.24

3 Layer MLP Results		

Feature	Accuracy (%)	Loss (%)

ReLU+Adam	98.27	8.78
ReLU+Batch+Adam	98	8.88
ReLU+Dropout+Adam	97.6	8.52

5 Layer MLP Results		

Feature	Accuracy (%)	Loss (%)

ReLU+Adam	97.65	11.87
ReLU+Batch+Adam	97.54	10.31
ReLU+Dropout+Adam	97.06	13.08

Observation:

1. For 2 hidden layers, the number of nodes taken in each layer are 364 and 72.
2. For 3 hidden layers, the number of nodes taken in each layer are 250,100,85.
3. For 5 hidden layers, the number of nodes taken in each layer are 235,102,65,50,35.
4. ReLU activation function with Adam optimizer is implemented with Batch normalization and dropout.
5. Upon seeing the results, for 3 MLP the Multi-class log loss is lesser(8.52) compared to 5 MLP(>10).
6. With respect to accuracy, again 3 MLP produces better results(98.27).